Moltbot on Your Copilot+ PC: A Beginner's Cookbook

No coding experience required. Just a Copilot+ PC and 30-45 minutes.

---

1. What is Moltbot? (And why you should be excited AND careful)

OpenClaw (formerly Moltbot, formerly Clawdbot) is an open-source AI personal assistant that runs locally on your computer. It was created by Austrian software engineer Peter Steinberger and released in late 2025. Within days, it became one of the fastest-growing open-source projects in GitHub history, surpassing 100,000 stars in its first two months.

The key difference between Moltbot and a chatbot like ChatGPT or Copilot: **a chatbot *tells* you things -- Moltbot *does* things.** It can read your files, write documents, run commands, install software, manage your calendar, and take autonomous actions across your apps. Scientific American described it as "AI with hands" that "follows almost any order" -- when given an objective, it "will break the objective into steps, find tools, install them, troubleshoot them and attempt to solve any obstacles."

Press coverage:

- **Scientific American:** "Moltbot Is an Open-Source AI Agent That Runs Your Computer" -- Called it "AI with hands" that can download, install, and run software autonomously. Noted it gives AI the ability to "run commands and manipulate files" while remembering user preferences.
- **Wired:** "The Viral AI Assistant Everyone's Talking About" -- Covered the explosive growth and community momentum. AI researcher Andrej Karpathy praised it publicly; investor Chamath Palihapitiya said Moltbot helped him save 15% on car insurance in minutes.
- **Platformer:** "Maybe someday you'll have a genie in your laptop. Today is not that day." -- A more skeptical review. Casey Newton found a gap between the hype (YouTube creators promising "24/7 AI employees") and reality. Her attempt to build a personalized morning briefing ultimately failed, leading her to uninstall it. An important reminder that this technology is still early.

The name: Originally "Clawdbot," Anthropic asked for a rename over trademark concerns with Claude. It became "Moltbot" (a nod to how lobsters molt their shells to grow), then "OpenClaw" in early 2026. (Read more on Wikipedia)

And here's the cool part: we're going to run it entirely on your laptop's AI chip. No cloud. No subscription. No data leaving your device. Ever.

BE CAREFUL -- This is powerful!

Moltbot is a REAL agent that can ACTUALLY modify your files and run commands on your computer. It's like giving someone the keys to your house. Security researchers -- including teams at Palo Alto Networks -- have warned that AI agents like this can "expose sensitive information and bypass security boundaries."

Things to know before you start:

- The AI running on your laptop is small (3.8 billion parameters -- for comparison, ChatGPT uses models 100x larger). **It WILL make mistakes sometimes.**
- **NEVER point it at important files** until you're comfortable with how it works.
- **Start in a safe test folder.** We'll set one up for you.
- If you tell it to delete something, **it will try to delete it.**
- Even Platformer's Casey Newton -- a seasoned tech journalist -- found the technology wasn't ready for complex real-world tasks yet. **Manage your expectations.** This is a proof of concept, not a finished product.

- Each action takes about 12 seconds. Be patient.

Power + responsibility. Start small, build confidence, then explore.

---

2. The brain behind the agent (and why "local" changes everything)

Here's something important to understand: **Moltbot isn't an AI model itself. It's the hands and feet -- the part that takes action.** It needs a separate AI model as its "brain" to decide what to do.

Think of it like a car: Moltbot is the steering wheel, pedals, and wheels. The AI model is the driver. You get to choose the driver.

The cloud option: world-class AI, but your data leaves the building

Moltbot is designed to work with many AI backends:

AI Model	Company	What you get
Claude Opus 4.5	Anthropic	One of the most capable AI models available today
GPT-4o / GPT-4.5	OpenAI	The model behind ChatGPT
Gemini 2.5 Pro	Google	Google's frontier model
Grok	xAI	Elon Musk's AI company
...and dozens more	Various	Via OpenRouter, you can access 200+ models

When you pair Moltbot with a frontier model like Claude Opus 4.5, the results can be remarkable. These models have hundreds of billions of parameters, trained on enormous datasets. They can reason through complex multi-step problems, write sophisticated code, and handle nuanced instructions with high reliability.

But there's a trade-off.

When Moltbot uses a cloud AI, every message you send -- including the contents of your files, your commands, your personal data -- travels over the internet to that company's servers. The AI reads it, processes it, and sends a response back.

That means:

- **Your data is on someone else's servers.** Even if companies promise not to misuse it, the data still physically leaves your device.
- **You're paying per use.** API costs add up -- heavy use of Claude Opus 4.5 can cost $50-100+/month easily.
- **You need an internet connection.** No wifi, no AI.
- **You're subject to rate limits.** Use too much too fast and you'll get throttled or cut off.
- **Privacy policies can change.** What's private today might be training data tomorrow.

For many tasks -- writing code, brainstorming, general research -- this is perfectly fine. But there are scenarios where sending data to the cloud is a non-starter.

The local option: your AI, your data, your rules

Now imagine the same AI agent, but the "brain" runs entirely on a chip inside your laptop. That's what we're doing with **Phi Silica on your Copilot+ PC's NPU**.

What is Phi Silica? It's a 3.8 billion parameter AI model made by Microsoft, specifically optimized to run on the Neural Processing Unit (NPU) inside Copilot+ PCs. It's the same model that powers Windows features like Recall, Click to Do, and Live Translation. It's already on your device -- you just need to activate it.

What's an NPU? It's a dedicated AI chip, separate from your CPU and GPU. It's designed to run AI workloads efficiently and with low power consumption. Think of it as a lane on the highway reserved just for AI -- it doesn't slow down everything else on your computer.

Here's the comparison:

	Cloud AI (e.g. Claude Opus)	Local AI (Phi Silica on NPU)
Intelligence	Exceptional (hundreds of billions of parameters)	Good for structured tasks (3.8B parameters)
Speed	2-5 seconds per response	~12 seconds per response
Privacy	Data sent to company servers	Data never leaves your device
Cost	$0.01-0.10+ per request	$0, forever
Internet required	Yes	No
Multi-step reasoning	Excellent	Limited (stateless -- each turn is independent)
Tool selection	Handles 20+ tools reliably	4 tools max (read, write, edit, exec)
Best for	Complex tasks, coding, analysis	File management, simple automation, demos

Why local matters for personal data

Consider what an AI agent might touch on your computer:

- **Financial documents** -- tax returns, bank statements, budgets
- **Medical records** -- prescriptions, lab results, insurance claims
- **Legal documents** -- contracts, NDAs, court filings
- **Personal communications** -- emails, messages, notes
- **Business data** -- client lists, proposals, internal memos
- **Photos and personal files** -- things meant only for you

With a cloud model, all of that data passes through someone else's infrastructure. With Phi Silica running locally, **none of it ever leaves your machine.** Not one byte. Not one packet. The AI processes everything right on the chip inside your laptop.

This isn't just a nice-to-have. For some people and industries, it's the difference between "we can use AI" and "we absolutely cannot use AI":

- **Healthcare workers** handling patient data (HIPAA)
- **Lawyers** with attorney-client privileged materials
- **Financial advisors** with client portfolios
- **Government employees** with classified or sensitive information
- **Anyone** who simply doesn't want their personal files on someone else's server

The honest trade-off

Local AI on a 3.8B model is not a replacement for Claude Opus or GPT-4. It's smaller, slower, and less capable at complex reasoning. It can't hold a conversation across turns (our implementation uses stateless mode for reliability). It sometimes picks the wrong tool or hallucinates parameters.

But it can reliably read files, write files, run commands, and edit documents -- entirely offline, entirely private, at zero cost. For many practical tasks, that's exactly what you need.

And this is just the beginning. As NPU hardware improves and local models get larger and smarter, the gap will shrink. Today's 3.8B model on a laptop NPU is a proof of concept. Tomorrow's 7B, 13B, or larger models on next-generation NPUs could rival cloud models for many everyday tasks -- while keeping your data exactly where it belongs: on your device.

---

3. What's cool about running it on YOUR Copilot+ PC?

Your Copilot+ PC has a special chip called an **NPU** (Neural Processing Unit). It's the same chip that powers Windows AI features you might already use:

- **Recall** (searching your PC history)
- **Click to Do** (smart actions on screen)
- **Live Translation**

We're using that **same AI chip** to run Moltbot. Here's why that matters:

Benefit	What it means for you
Total privacy	Your data NEVER leaves your device. Not one byte.
Zero cost	No subscription, no API fees, no token counting. Run it forever.
Works offline	On a plane, no wifi, in a coffee shop -- doesn't matter.
Already installed	Phi Silica (the AI model) is already on your Copilot+ PC.

It's YOUR AI on YOUR hardware.

---

4. What can it do? (Simple examples)

Once set up, you'll type commands like:

> "*Read my notes.txt file and tell me what's in it*"

> "*Create a new file called shopping-list.txt with milk, eggs, and bread*"

> "*What time is it?*" *(runs a system command to find out)*

> "*Edit my file and change 'Monday' to 'Tuesday'*"

Each action takes about **12 seconds** on the NPU. That's the AI thinking -- be patient with it!

---

5. What do I need?

Checklist before you start:

- ☐ A **Copilot+ PC** (Surface Laptop 7 Intel, Surface Pro 11, ASUS, Dell, HP, Lenovo, Samsung -- any Copilot+ PC)
- ☐ **Windows 11** version 24H2 or later
- ☐ About **30-45 minutes** of time
- ☐ Willingness to **copy and paste** some commands

A note on devices: This guide is tested on **Intel Lunar Lake** (Surface Laptop 7). Qualcomm Snapdragon X support is coming next -- we started with Intel simply because that's the device we had open when we started this project! No favorites here.

Yes, you'll need to install a few things. Don't worry -- it's all just clicking "Next" a few times and copy-pasting commands. No actual coding required.

---

6. Let's set it up!

Step 6.1: Install VS Code

What it is: A free text editor from Microsoft. We need it to activate the AI on your device.

1. Go to **https://code.visualstudio.com/**
2. Click **"Download for Windows"**
3. Run the installer
4. Accept all defaults -- click **Next** until it's done
5. Open VS Code when finished

> ***Checkpoint:*** *You see VS Code open with a Welcome tab.*

Step 6.2: Install the AI Toolkit extension

What it is: A VS Code add-on that manages AI models on your device.

1. In VS Code, click the **Extensions icon** on the left sidebar (looks like 4 squares)
2. In the search box, type: **AI Toolkit**
3. Find **"AI Toolkit"** by Microsoft (has a blue icon)
4. Click **Install**
5. Wait for it to finish

> ***Checkpoint:*** *A new AI Toolkit icon appears in the left sidebar (looks like a sparkle/star).*

Step 6.3: Load Phi Silica from Windows AI API

This is the key step -- we're activating the AI that's already on your device!

1. Click the **AI Toolkit icon** in the left sidebar
2. Click on **"Model Catalog"**
3. You'll see tabs at the top: **Foundry Local | Ollama | Windows AI API**
4. Click on the **"Windows AI API"** tab
5. Find **"Phi Silica"** in the list
6. Click **"Add"** to load it
7. Wait for it to initialize (you might see a loading indicator)

> ***Checkpoint:*** *Phi Silica appears in your "My Models" or loaded models section.*

Step 6.4: Test Phi Silica in VS Code

Before we go further, let's make sure your NPU AI is actually running.

1. In AI Toolkit, click on **Phi Silica**
2. Click **"Try in Playground"** or **"Open Playground"**
3. Type a simple message: **"Hello, what can you do?"**
4. Wait 10-15 seconds for a response

> ***Checkpoint:*** *You get a response! Your NPU is working.*
>
> ***If this doesn't work, STOP here.*** *There's no point continuing until the AI model responds in the playground. Check that your Windows is up to date (24H2+) and that your device is a Copilot+ PC.*

Step 6.5: Install Node.js

What it is: The engine that runs Moltbot. Think of it like how you need a DVD player to play a DVD -- Moltbot needs Node.js to run.

1. Go to **https://nodejs.org/**
2. Click the big green **LTS** button (LTS means "stable version")
3. Run the installer
4. Click **Next**, accept the license, **Next**, **Next**, **Install**
5. If you see a checkbox saying **"Automatically install necessary tools"** -- check it
6. Click **Finish**

IMPORTANT: Close ALL PowerShell or Command Prompt windows you have open, then reopen them. Node.js needs a fresh window to be recognized.

> ***Checkpoint:*** *Open PowerShell (press the Windows key, type **PowerShell**, click it) and type:*
>
> ```
> node --version
> ```
>
> *You should see something like* `v24.13.0`*. If you see "not recognized", close PowerShell completely and reopen it.*

Step 6.6: Allow PowerShell scripts

PowerShell blocks scripts by default for security. We need to relax that slightly.

Copy and paste this into PowerShell:

```
Set-ExecutionPolicy -ExecutionPolicy RemoteSigned -Scope CurrentUser
```

When asked, type **A** (for "Yes to All") and press Enter.

Step 6.7: Install OpenClaw (Moltbot)

Copy and paste this into PowerShell:

```
npm install -g openclaw
```

Wait for it to finish (1-2 minutes). You'll see a lot of text scrolling by -- that's normal.

> ***Checkpoint:*** *You see "added X packages" at the end with no red error messages.*

Step 6.8: Get the NPU patch from GitHub

What this does: Downloads the code that teaches OpenClaw how to use your NPU.

Copy and paste these commands **one at a time**, pressing Enter after each:

```
cd $HOME
```

```
git clone https://github.com/frankcx1/pi-mono.git pi-mono-npu
```

```
cd pi-mono-npu
```

```
git checkout foundry-local-npu-support
```

If you see "git is not recognized":

1. Go to **https://git-scm.com/downloads/win**
2. Download and install (accept all defaults)
3. **Close and reopen PowerShell**
4. Try the commands above again

> ***Checkpoint:*** *The last command says "Switched to branch 'foundry-local-npu-support'" or "Already on 'foundry-local-npu-support'".*

Step 6.9: Build the NPU patch

What this does: Compiles the source code into something OpenClaw can use. Think of it like converting a recipe into an actual dish.

Copy and paste these commands one at a time:

```
npm install
```

This downloads the ingredients (takes 2-5 minutes, lots of text is normal).

```
npm run build
```

This does the actual cooking (1-2 minutes).

> ***Checkpoint:*** *No red error messages. The last few lines mention packages building successfully.*

Step 6.10: Apply the NPU patch to OpenClaw

What this does: Copies the NPU-aware code into your OpenClaw installation.

Copy and paste this entire command (it's one long line):

```
Copy-Item -Path ".\packages\ai\dist\providers\openai-completions.js" -Destination
"$env:APPDATA\npm\node_modules\openclaw\node_modules\@mariozechner\pi-
ai\dist\providers\openai-completions.js" -Force
```

> ***Checkpoint:*** *No output means success! (In PowerShell, silence is golden.)*

Step 6.11: Set up OpenClaw

Copy and paste this into PowerShell:

```
npx openclaw setup
```

Follow any prompts. When asked about API keys, you can skip them -- we're running locally!

Step 6.12: Create the Foundry Local configuration

We need to tell OpenClaw to use Phi Silica on your NPU. Copy and paste this entire block:

```
$configDir = "$env:USERPROFILE\.openclaw"
if (!(Test-Path $configDir)) { New-Item -ItemType Directory -Path $configDir -Force }

@'
{
  "models": {
    "providers": {
      "foundry-local": {
        "baseUrl": "http://localhost:5272/v1",
        "apiKey": "not-needed",
        "api": "openai-completions",
        "models": [{ "id": "phi-silica", "name": "Phi Silica (NPU)" }]
      }
    }
  },
  "agents": {
    "defaults": {
      "model": { "primary": "foundry-local/phi-silica" }
    }
  }
}
'@ | Out-File -FilePath "$configDir\openclaw.json" -Encoding UTF8
```

> ***Checkpoint:*** *No errors. You can verify by running:*
>
> ```
> Get-Content "$env:USERPROFILE\.openclaw\openclaw.json"
> ```
>
> *You should see the JSON configuration we just wrote.*

Step 6.13: Create a safe test folder

Before we start, let's make a safe place to experiment where you won't accidentally touch important files:

```
$testDir = "$env:USERPROFILE\.openclaw\workspace"
if (!(Test-Path $testDir)) { New-Item -ItemType Directory -Path $testDir -Force }
"Hello from my Copilot+ PC!" | Out-File -FilePath "$testDir\hello.txt" -Encoding UTF8
```

Step 6.14: Start chatting!

Make sure Phi Silica is loaded in AI Toolkit (Step 6.3-6.4). Then run:

```
npx openclaw agent --local --session-id my-first-test --message "read
$env:USERPROFILE\.openclaw\workspace\hello.txt"
```

Wait about 12 seconds...

You should see the AI read your file and tell you what's in it!

---

7. Try these commands!

Now that it's working, try these one at a time. Each takes about 12 seconds.

Say hello:

```
npx openclaw agent --local --session-id test --message "hello, what can you help me with?"
```

Create a file:

```
npx openclaw agent --local --session-id test --message "use the write tool to create a file at
C:\Users\$env:USERNAME\.openclaw\workspace\shopping.txt with content: milk, eggs, bread,
butter"
```

Read it back:

```
npx openclaw agent --local --session-id test --message "read
C:\Users\$env:USERNAME\.openclaw\workspace\shopping.txt"
```

Run a command:

```
npx openclaw agent --local --session-id test --message "run the command: powershell -c Get-
Date"
```

Write a haiku:

```
npx openclaw agent --local --session-id test --message "write a haiku about computers"
```

---

8. Troubleshooting (in plain English)

"'node' is not recognized" Close PowerShell completely, reopen it, try again. Node.js needs a fresh window after installation.

"'git' is not recognized" Install Git from https://git-scm.com/downloads/win, then close and reopen PowerShell.

"npm ERR!" with red text during install Try running PowerShell as Administrator (right-click PowerShell, "Run as administrator") and try the npm install command again.

No output or blank responses Make sure Phi Silica is loaded in AI Toolkit. Go back to Step 6.4 and verify it responds in the playground.

Takes more than 30 seconds with no response The AI might have unloaded from memory. Open AI Toolkit in VS Code, click on Phi Silica, make sure it's active. Try again.

"file not found" errors Check your username in the path. Open File Explorer and go to `C:\Users\` to see your exact username. It's case-sensitive!

"npm run build" fails with errors Make sure you ran `npm install` first and it completed without errors. If it still fails, try deleting the `node_modules` folder and running `npm install` again:

```
Remove-Item -Recurse -Force node_modules
npm install
npm run build
```

---

9. Stuck? Ask AI for help!

If you get stuck at any point, copy this into **Microsoft Copilot**, **Claude**, or **ChatGPT**:

Copy this entire block:

```
I'm following a beginner's guide to set up Moltbot/OpenClaw on my Copilot+ PC
with Phi Silica NPU.

My setup:
- Device: [Surface Laptop 7 / Surface Pro 11 / other - fill in yours]
- Windows version: [run 'winver' to check]
- Current step I'm stuck on: [describe which step number]

The error or problem I'm seeing:
[paste any error message here]

What I've tried:
[describe what you attempted]

Can you help me troubleshoot this specific issue?
```

Pro tip: Screenshots help! Press **Win + Shift + S** to capture your screen, then paste it right into the AI chat.

The AI assistants are great at debugging setup issues -- that's literally how this whole project was built!

---

10. What's next?

- **Share what you built!** Tag **#Moltbot #CopilotPlus #LocalAI** on social media
- **Follow the project:** https://github.com/badlogic/pi-mono
- **Snapdragon X** support coming soon!
- The more people try this, the better it gets

---

11. The story behind this

This started as a weekend experiment -- "I wonder if I could get an AI agent running entirely on my Surface Laptop's NPU?" A few hours of vibe coding with Claude Code later, it actually worked.

The whole project was built in one session -- from "is this even possible?" to working code with documentation. If a random person with a laptop can do this, imagine what the community can build together.

If you try this and build something cool, share it! That's how open source works.